

## Importing the required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

## Loading the dataset

```
file_path = "./house_prices_dataset.csv"
df = pd.read_csv(file_path)
df
```

	Location	Size_sqft	Bedrooms	Bathrooms	House_Age	Garage
Pool \						
0	Chicago	1395	3	1	26	1
0						
1	Houston	1528	4	2	28	1
0						
2	Los Angeles	2165	4	4	46	1
0						
3	Houston	1498	5	3	27	1
0						
4	Houston	1195	5	4	34	1
1						
..	...	...	...	...	...	...
...						
495	New York	4475	3	4	14	1
1						
496	Los Angeles	3854	3	4	16	1
1						
497	New York	2943	1	1	39	1
0						
498	New York	2111	1	4	17	0
0						
499	New York	3991	4	2	8	1
0						
	Distance_to_City_Center_miles					Price
0		12.077516				7.257565e+05
1		22.750817				7.264075e+05
2		27.593580				1.050788e+06

```

3          28.552394  7.240178e+05
4          17.525116  7.342087e+05
..          ...
495        22.624345  1.917976e+06
496        27.292740  1.643563e+06
497        22.882102  1.270751e+06
498        18.116082  1.127655e+06
499        19.796916  1.723863e+06

```

```
[500 rows x 9 columns]
```

## Preprocessing and Handling the Missing Values

```
# dropping the duplicate rows and any missing values
```

```
df.drop_duplicates(inplace=True)
```

```
df.dropna(inplace=True)
```

```
df
```

	Location	Size_sqft	Bedrooms	Bathrooms	House_Age	Garage
Pool \						
0	Chicago	1395	3	1	26	1
0						
1	Houston	1528	4	2	28	1
0						
2	Los Angeles	2165	4	4	46	1
0						
3	Houston	1498	5	3	27	1
0						
4	Houston	1195	5	4	34	1
1						
..	...	...	...	...	...	...
...						
495	New York	4475	3	4	14	1
1						
496	Los Angeles	3854	3	4	16	1
1						
497	New York	2943	1	1	39	1
0						
498	New York	2111	1	4	17	0
0						
499	New York	3991	4	2	8	1
0						
	Distance_to_City_Center_miles				Price	
0		12.077516			7.257565e+05	
1		22.750817			7.264075e+05	
2		27.593580			1.050788e+06	
3		28.552394			7.240178e+05	
4		17.525116			7.342087e+05	

```

..
495          22.624345  1.917976e+06
496          27.292740  1.643563e+06
497          22.882102  1.270751e+06
498          18.116082  1.127655e+06
499          19.796916  1.723863e+06

```

[500 rows x 9 columns]

## Encoding the categorical variable (Location)

```

# encoding the categorical variables using OneHotEncoder
encoder = OneHotEncoder(drop='first', sparse_output=False)
encoded_location = encoder.fit_transform(df[['Location']])
# converting the encoded location into dataframe with proper columns
encoded_location_df = pd.DataFrame(encoded_location, columns =
encoder.get_feature_names_out(['Location']))
# adding the encoded location (categorical feature) to the dataframe
df = df.drop( columns= ['Location'])
df = pd.concat([df, encoded_location_df], axis = 1)
df

```

	Size_sqft	Bedrooms	Bathrooms	House_Age	Garage	Pool	\
0	1395	3	1	26	1	0	
1	1528	4	2	28	1	0	
2	2165	4	4	46	1	0	
3	1498	5	3	27	1	0	
4	1195	5	4	34	1	1	
..	...	...	...	...	...	...	
495	4475	3	4	14	1	1	
496	3854	3	4	16	1	1	
497	2943	1	1	39	1	0	
498	2111	1	4	17	0	0	
499	3991	4	2	8	1	0	

	Distance_to_City_Center_miles	Price	Location_Houston	\
0	12.077516	7.257565e+05	0.0	
1	22.750817	7.264075e+05	1.0	
2	27.593580	1.050788e+06	0.0	
3	28.552394	7.240178e+05	1.0	
4	17.525116	7.342087e+05	1.0	
..	...	...	...	
495	22.624345	1.917976e+06	0.0	
496	27.292740	1.643563e+06	0.0	
497	22.882102	1.270751e+06	0.0	
498	18.116082	1.127655e+06	0.0	
499	19.796916	1.723863e+06	0.0	

	Location_Los Angeles	Location_New York	Location_San Francisco
0	0.0	0.0	0.0

1	0.0	0.0	0.0
2	1.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
...	...	...	...
495	0.0	1.0	0.0
496	1.0	0.0	0.0
497	0.0	1.0	0.0
498	0.0	1.0	0.0
499	0.0	1.0	0.0

[500 rows x 12 columns]

## Scaling the data

```
# scaling the data using StandardScaler
X = df.drop(columns=['Price'])
y = df['Price']
X_scaled = StandardScaler().fit_transform(X)
```

## Splitting into Train and Test data

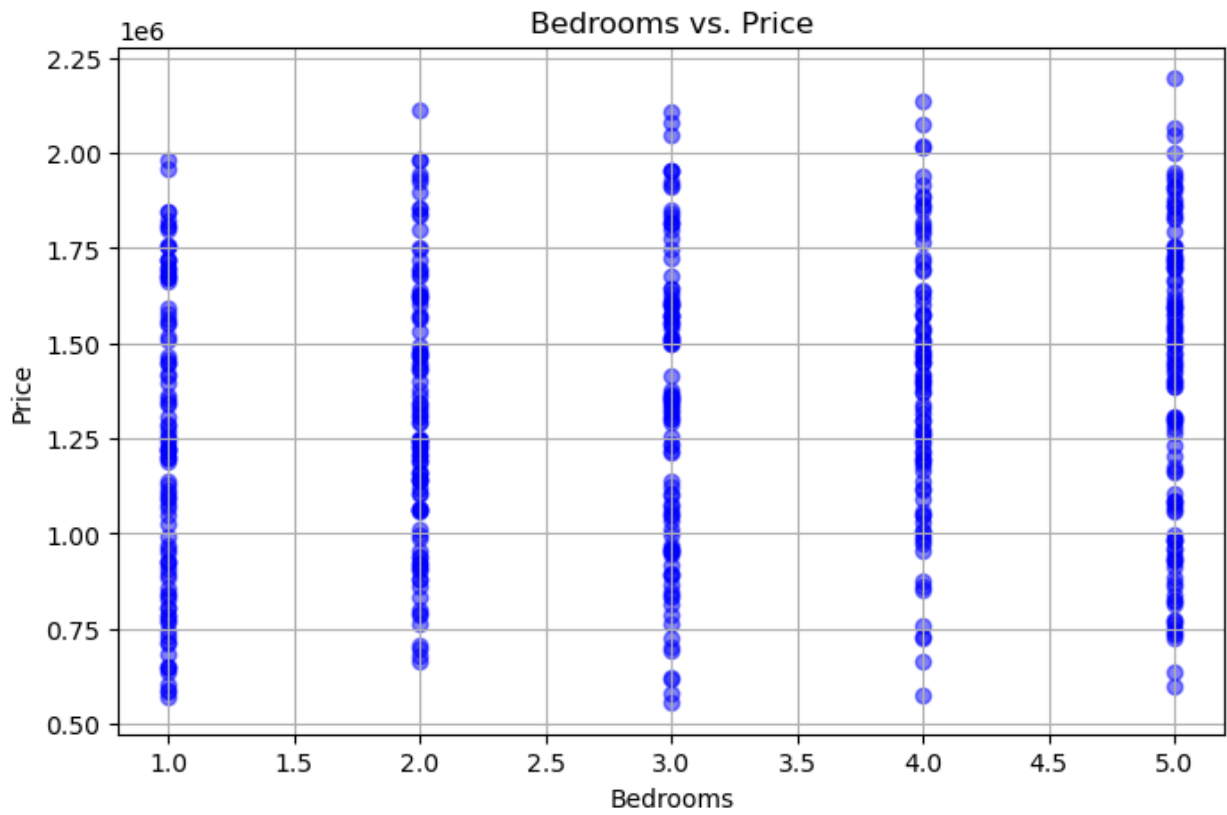
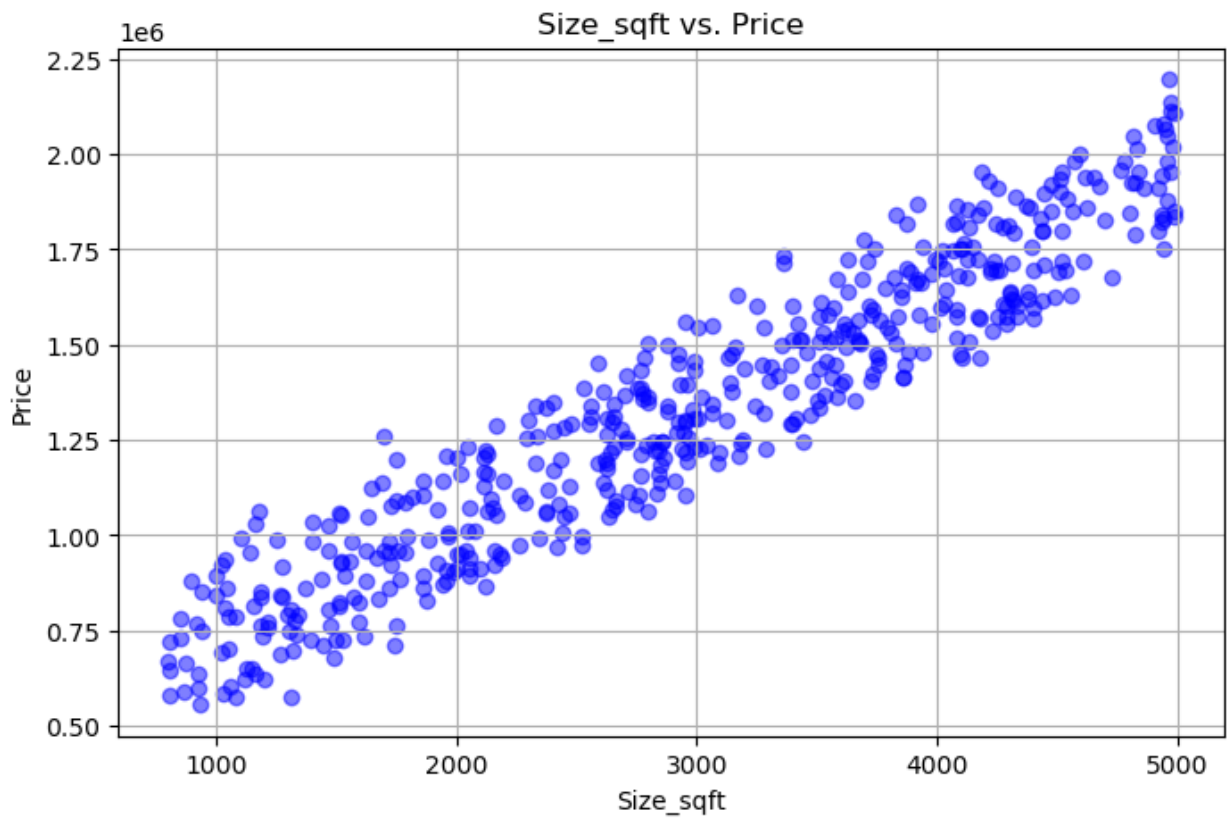
```
# splitting the data into train and test in 80 / 20 ratio
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)
```

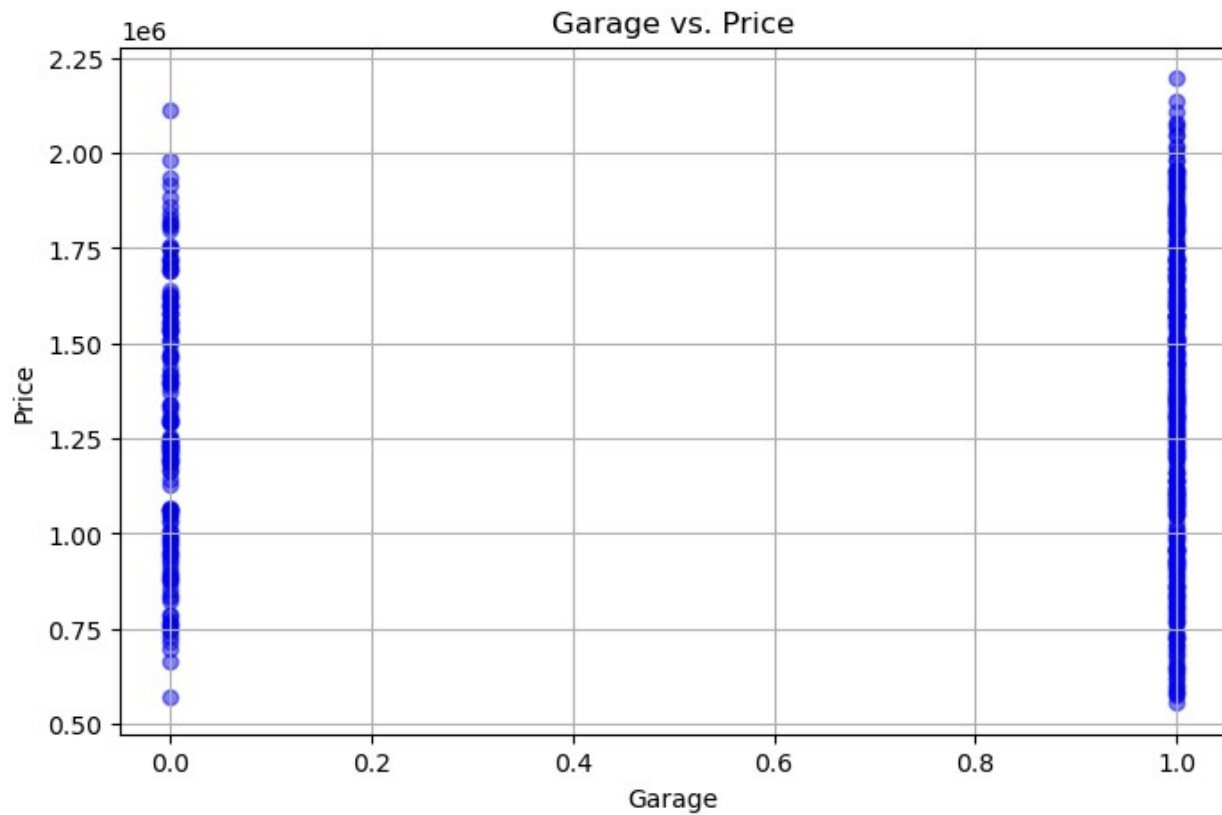
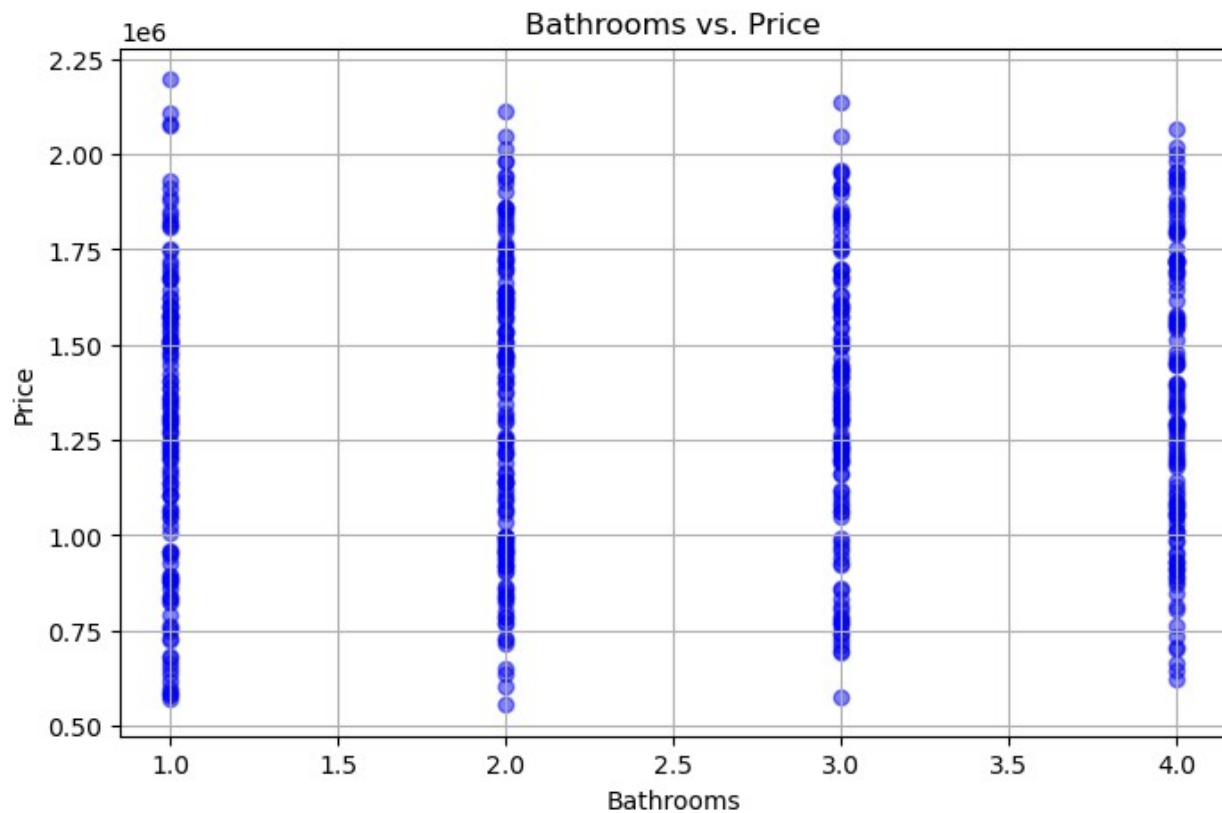
## Exploratory data analysis


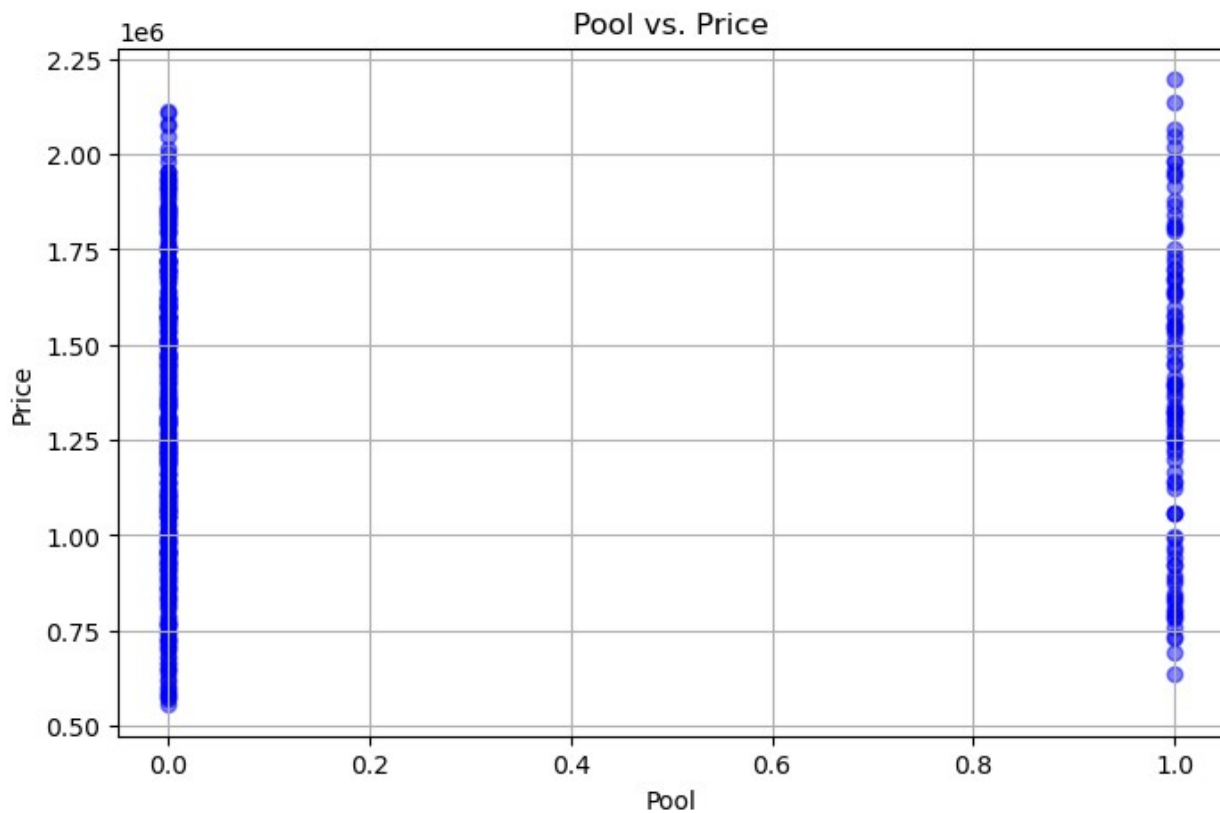
### 1.scatter plots between features and the house prices

```
# using scatter plot to understand the trends
features = ['Size_sqft', 'Bedrooms', 'Bathrooms', 'Garage', 'Pool',
'Distance_to_City_Center_miles']

for col in features:
    plt.figure(figsize=(8, 5))
    plt.scatter(df[col], df['Price'], alpha=0.5, color='blue')
    plt.title(f"{col} vs. Price")
    plt.xlabel(col)
    plt.ylabel("Price")
    plt.grid(True)
    plt.show()
```








## 2. Histogram of the house prices

```
plt.figure(figsize=(10, 5))
plt.hist(df['Price'], bins=40, color='blue', edgecolor='black',
alpha=0.5)
plt.title("Distribution of House Prices")
plt.xlabel("Price")
plt.ylabel("Frequency")
plt.show()
```



## 3. Info about various statical measures of the given dataset

```
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 500 entries, 0 to 499
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	Size_sqft	500 non-null	int64
1	Bedrooms	500 non-null	int64
2	Bathrooms	500 non-null	int64
3	House_Age	500 non-null	int64
4	Garage	500 non-null	int64
5	Pool	500 non-null	int64
6	Distance_to_City_Center_miles	500 non-null	float64
7	Price	500 non-null	float64
8	Location_Houston	500 non-null	float64



9	Location_Los Angeles	500 non-null	float64
10	Location_New York	500 non-null	float64
11	Location_San Francisco	500 non-null	float64

dtypes: float64(6), int64(6)

memory usage: 47.0 KB

	Size_sqft	Bedrooms	Bathrooms	House_Age	Garage \
count	500.000000	500.000000	500.000000	500.000000	500.000000
mean	2971.950000	3.008000	2.458000	23.458000	0.702000
std	1169.733097	1.428291	1.114774	14.209435	0.457838
min	803.000000	1.000000	1.000000	0.000000	0.000000
25%	2005.000000	2.000000	1.000000	11.000000	0.000000
50%	2956.000000	3.000000	2.000000	23.000000	1.000000
75%	3980.500000	4.000000	3.000000	35.000000	1.000000
max	4988.000000	5.000000	4.000000	49.000000	1.000000

	Pool	Distance_to_City_Center_miles	Price \
count	500.000000	500.000000	5.000000e+02
mean	0.194000	15.200508	1.329449e+06
std	0.395825	8.373759	3.780628e+05
min	0.000000	0.500343	5.558926e+05
25%	0.000000	7.903742	1.029014e+06
50%	0.000000	15.467201	1.336224e+06
75%	0.000000	22.237744	1.619568e+06
max	1.000000	29.916879	2.196013e+06

	Location_Houston	Location_Los Angeles	Location_New York \
count	500.000000	500.000000	500.000000
mean	0.186000	0.182000	0.218000
std	0.389496	0.386231	0.413301
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	Location_San Francisco
count	500.000000
mean	0.190000
std	0.392694
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

# Training the data using different models

## 1.Linear Regression

```
# using linear regression to train
model1 = LinearRegression()
model1.fit(X_train, y_train)
y_pred1 = model1.predict(X_test)

# performance metrics of linear regression
print("RMSE_LINEAR_REGRESSION : ", np.sqrt(mean_squared_error(y_pred1,
y_test)))
print("MAE_LINEAR_REGRESSION : ", mean_absolute_error(y_pred1,
y_test))
print("R2_SCORE_LINEAR_REGRESSION : ", r2_score(y_pred1, y_test))

RMSE_LINEAR_REGRESSION :  11719.545666352768
MAE_LINEAR_REGRESSION :  10334.125078503075
R2_SCORE_LINEAR_REGRESSION :  0.998960835812596
```

## 2.Neural Networks

```
# using multi layer feed forward network to train
model2 = MLPRegressor(hidden_layer_sizes= (100,
50),activation='relu' ,solver='adam', max_iter= 500, random_state= 42)
model2.fit(X_train, y_train)
y_pred2 = model2.predict(X_test)

# calculating performance metrics of neural networks
print("RMSE_NEURAL_NETWORKS : ", np.sqrt(mean_squared_error(y_pred2,
y_test)))
print("MAE_NEURAL_NETWORKS : ", mean_absolute_error(y_pred2, y_test))
print("R2_SCORE_NEURAL_NETWORKS : ", r2_score(y_pred2, y_test))

RMSE_NEURAL_NETWORKS :  1369368.162748587
MAE_NEURAL_NETWORKS :  1323044.0458968843
R2_SCORE_NEURAL_NETWORKS :  -8266.788117091606

/usr/lib/python3/dist-packages/sklearn/neural_network/
_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (500) reached and the optimization
hasn't converged yet.
  warnings.warn(
```

## 3.Decision Tree

```
# using decision tree to train
model3 = DecisionTreeRegressor()
model3.fit(X_train, y_train)
```

```
y_pred3 = model3.predict(X_test)

# calculating performance metrics of decision tree
print("RMSE_DECISION_TREE : ", np.sqrt(mean_squared_error(y_pred3,
y_test)))
print("MAE_DECISION_TREE : ", mean_absolute_error(y_pred3, y_test))
print("R2_SCORE_DECISION_TREE : ", r2_score(y_pred3, y_test))

RMSE_DECISION_TREE :  91841.89581855651
MAE_DECISION_TREE :  74880.42626300002
R2_SCORE_DECISION_TREE :  0.9390259587806761
```